

Rohit Khazanchi

Vancouver, BC — 647-507-0212 — rohitkhazanchi121@gmail.com — linkedin.com/in/rohitkhazanchi

Summary

Senior Data Engineer with 8+ years of experience building large-scale data platforms on Azure across mining, retail, and SaaS domains. Strong expertise in streaming and batch processing, data governance, and operational reliability, with a focus on building trusted, scalable data systems.

Experience

Senior Data Developer — Teck Resources Ltd

Dec 2020 – Present

Vancouver, BC

- Led the development of an enterprise plant optimization platform across multiple mine sites, defining technical roadmap, data contracts (schema, latency, availability) and supporting the hiring of key roles.
- Architected and implemented scalable streaming pipelines processing 500M+ IoT sensor events using Azure Event Hubs; transformed and aggregated data into time-series storage for analytics and product insights.
- Designed and built incremental batch pipelines using ADF, Databricks, and ADLS Gen2 to ingest data from files, databases, and APIs into curated analytical datasets.
- Defined and enforced data contracts across upstream systems, improving data reliability and reducing upstream-related incidents by approximately 75%.
- Built end-to-end observability using custom metrics, structured logging, distributed tracing, KQL queries, and Grafana dashboards (volume, latency and data quality pattern), reducing mean time to resolution (MTTR) by more than 50%.
- Implemented data governance and compliance frameworks with fine-grained RBAC, ensuring alignment with enterprise data privacy and regulatory standards.
- Migrated pipelines from row-based to bulk processing patterns, achieving a 50% throughput improvement without increasing infrastructure cost.
- Designed self-healing pipelines with automated anomaly detection and retry mechanisms, improving SLA adherence for critical data workflows.
- Built and maintained a data glossary and lineage documentation, enabling data scientists, analysts, and compliance teams to understand the ownership and dependencies of the data set.
- Developed a RAG-based LLM assistant using LangGraph and Postgres Vector, enabling 50+ users to retrieve technical documentation 60% faster via a grounded chat interface.
- Mentored two junior engineers through structured feedback and code reviews, enabling them to independently own projects within a year.

Module Lead, Data Engineering — IGT Solutions

May 2017 – Dec 2020

Gurugram, India

- Led a team of three engineers to design and deliver enterprise-grade ETL pipelines using Azure Data Factory, integrating data from relational databases and APIs into analytics-ready datasets.
- Optimized high-frequency T-SQL queries by analyzing execution plans, restructuring indexes, and refactoring stored procedures, improving performance by up to 60% for reporting workloads.
- Built and optimized data models and Power BI reporting layers, enabling business stakeholders to derive actionable insights from integrated datasets.
- Established data quality validation and monitoring processes to ensure completeness, consistency, and reliability across data pipelines.
- Collaborated directly with clients on complex debugging scenarios, reducing recurring incidents through root cause analysis instead of short-term fixes.

BI Specialist — Adobe

May 2015 – May 2017

Noida, India

- Designed scalable database schemas and data models for analytics across large-scale datasets used by multiple business teams.
- Optimized ETL processes to load millions of incremental records every 45 minutes, improving data freshness and downstream usability.
- Built consolidated dashboards for customer retention teams, centralizing multi-source data to identify upsell opportunities.

Previous Experience

Senior Systems Engineer — Infosys

Feb 2012 – May 2015

Projects

Plant Optimization Platform

- Instrumented streaming jobs with custom latency metrics integrated into Grafana dashboards, enabling proactive monitoring and reducing reactive incident handling.
- Built stream processing logic to normalize heterogeneous IoT sensor data across varying schemas, sampling rates, and units into a unified feature store.
- Designed alerting mechanisms for missing sensor data and inference failures, reducing detection time from hours to minutes while maintaining 99% ML availability.

Operational Data Warehouse & Reporting

- Built an end-to-end Azure SQL data warehouse using star schema design, incremental loading, and index-optimized analytical queries.
- Designed a tabular semantic layer with optimized DAX measures powering Power BI dashboards for business stakeholders.

Technical Skills

Languages: Python, PySpark, SQL (T-SQL, PostgreSQL), PL/SQL, KQL, Bash

Azure Data Platform: ADF, ADX (Kusto), ADLS Gen2, Event Hubs, Synapse

Distributed Systems: Apache Spark, Databricks, Delta Lake, Kafka, Airflow

Databases: PostgreSQL, TimescaleDB, SQL Server, Azure SQL, Cosmos DB, Oracle

Infrastructure: Kubernetes (AKS), Docker, Terraform, Azure DevOps, GitHub Actions

Observability: Grafana, logging, tracing, metrics, alerting pipelines

Education

MBA — NMIMS, Mumbai

2015 – 2017

B.Tech, Computer Science — Maharishi Dayanand University

2007 – 2011